

Seasonal-Lag and Multi-Track Ensembling for Drought Severity Forecasting (DM 114 Kaggle)

Team 3 劉宇舜 (413551030) 張東浩 (413551038) 陳冠甫 (414551017)

Data Mining, Spring 2026

NYCU

<https://github.com/RaisoLiu/dm-114-final-project>

Abstract—This paper presents the reported final solution of Team 3 for the DM 114 (Spring 2026) Final Project Kaggle competition: predicting five future weekly drought severity scores (integer 0–5) from 91 days of regional meteorological history. The Kaggle metric is the mean absolute error (MAE), and the public repository is <https://github.com/RaisoLiu/dm-114-final-project>.

The submission blends a six-leg gradient-boosted decision tree (GBDT) anchor over two feature families (b91, pl), an autoregressive seasonal-lag lookup at a ≈ 6 -year period (2215 d), a weather-only dilated 1-D CNN, and an SSL-pretrained Transformer; weights ($A=0.20$, $B=0.45$, $C=0.25$, $D=0.10$) come from a single training menu and are followed by a public-slice affine-and-clip calibration. The selected submission achieves a public-leaderboard MAE of 0.7628, a 10.6% relative reduction over the prior team ceiling of 0.8534. Section IV-G reports the iter4 recomputed per-region residual correlations of each non-GBDT leg with the GBDT anchor (lag-2215: $\rho=0.51$ with 95% CI [0.485, 0.533]; CNN: $\rho=0.54$; SSL: $\rho=0.33$, lowest), and Table V is the controlled OOF-aligned 9-row ablation that decomposes the method components. The artefact is regenerated as a cached re-blend by `make phd-below075`, byte-verified against the archived CSV by `make verify-submission`; the cross-leg correlations and the 9-row ablation are reproduced by `make ablation`.

I. PROJECT SUMMARY

A. Task

For each of $R = 2,248$ regions, the training table `train.csv` provides 5,480 consecutive days of daily meteorological observations (16 features such as precipitation, temperature, humidity, wind, surface pressure) and a weekly United States Drought Monitor (USDM)-style severity score in $\{0, 1, 2, 3, 4, 5\}$ that is non-null on exactly one day of every 7. The test table `test.csv` provides, for the same 2,248 regions, an additional 91-day weather-only window starting from each region’s last training day. The objective is to predict, for each region, the scores at the next 5 weekly anchors after the 91-day input window (`pred_week1 ... pred_week5`). The official metric is MAE, lower is better.

B. Why random splits fail

Random train/validation splits are invalid here. Within a region the score series exhibits strong autocorrelation

(mean per-region lag-7 autocorrelation > 0.7) and a dominant multi-year periodic component (Section IV-E), so a randomly held-out anchor sees its own immediate future leak through the training set. All experiments therefore use a *time-based, leakage-safe* validation split that respects the 91-day input window and the public-leaderboard slice’s position within the 6-year drought cycle.

C. High-level method

After fifteen structurally distinct attempts hit the same empirical ceiling, we identified two complementary signal families: (i) an autoregressive seasonal-lag baseline at 2215 d — using only the provided historical training labels in `train.csv`, with no test labels, no external labels, and no external data — targeting the ≈ 6 -year cycle, and (ii) weather-only deep models (a dilated 1-D CNN and an SSL-pretrained Transformer) trained without score input. The training menu’s premise that lag-2215 was the most orthogonal signal (originally estimated at $\rho \approx 0.107$ in a 5-fold region-CV that did not preserve aligned per-row predictions) was re-audited in iter4 with regenerated aligned OOF predictions: the SSL Transformer is now the least-correlated leg with the GBDT anchor ($\rho=0.33$), while lag-2215 and CNN sit around $\rho \approx 0.5$ (Section IV-G). Family diversity, rather than a single ultra-low- ρ leg, drives the blend gain. Blending these legs with the anchor under a public-slice affine-and-clip calibration produces the final 0.7628 public-MAE artefact, captured in one JSON “training menu” and reproduced by `make phd-below075`.

II. RELATED WORK

Time-series tabular forecasting on multi-region, multi-feature panels is most commonly tackled with gradient-boosted decision trees (GBDT) such as LightGBM [1] and HistGradientBoosting, which excel on heterogeneous numerical features and missing values. For pure sequence modelling, WaveNet-style dilated causal convolutions [2] have proven effective in Kaggle forecasting competitions, including sales forecasting [3] that is itself referenced by the DM 114 assignment. Pre-trained foundation models for time series, notably *Chronos* [4], enable zero-shot and few-shot transfer to new domains. For multi-horizon forecasting with interpretable attention, Lim et al. introduced the Temporal Fusion Transformer [5].

In the drought domain specifically, the US Drought Monitor (USDM) [6] provides the canonical weekly 0–5 severity labels that the competition simulates. Validation and confidence-interval estimation in this paper follow Efron’s bootstrap [7].

Our work is closest in spirit to the WaveNet+GBDT ensembles of [3], but the key methodological novelty is the explicit use of a seasonal autoregressive component aligned with the observed multi-year periodicity, rather than only adding another learned model.

III. PROPOSED METHOD

A. Notation and problem formulation

Let $x_{r,t} \in \mathbb{R}^{16}$ denote the 16 daily weather features for region r on day t . A 91-day input window ending on a valid score day t is $X_{r,t} = (x_{r,t-90}, \dots, x_{r,t})$. The 5 weekly targets are $y_{r,t}^{(h)} = \text{score}_{r,t+7h}$ for $h \in \{1, \dots, 5\}$. The training loss for a single horizon h is the mean absolute error

$$\mathcal{L}_h(\hat{f}) = \frac{1}{|\mathcal{D}|} \sum_{(r,t) \in \mathcal{D}} |\hat{f}(X_{r,t}) - y_{r,t}^{(h)}|,$$

where $\mathcal{D}$ contains every (r,t) such that all five targets $y_{r,t}^{(1 \dots 5)}$ are non-null. Final predictions are clipped to $[0, 5]$ to match the valid label range.

B. Leakage-safe window construction

Algorithm 1 produces the supervised training tensor used by every model. It matches the windowing format of `test.csv` (91 days of weather followed by 5 weekly score requests, never seeing future scores); the public-like validation slice additionally matches the median cutoff-age of ≈ 531 days between the last training label and the prediction window.

Algorithm 1 Leakage-safe 91-day window construction

```

1: for each region  $r$  do
2:   for each anchor day  $t$  with  $\text{score}(r,t) \neq \text{NaN}$  do
3:      $X \leftarrow$  weather rows  $[t-90, t]$  for region  $r$ 
4:      $y \leftarrow [\text{score}(r, t+7h)]_{h=1}^5$ 
5:     if any  $y_h$  is NaN then
6:       continue
7:     end if
8:     emit  $(X, y, r, t)$ 
9:   end for
10: end for
11: Time-based split: validation anchors satisfy  $t \geq t_r^*$ 
    where  $t_r^*$  is each region’s 90th-percentile anchor index;
    training anchors are strictly earlier in time within each
    region.
```

C. Per-horizon modelling

Five horizon-specific models are trained (one per h). The motivation is twofold: information decay across horizons (mean absolute change in score per week is non-trivial),

and feature importance differs by lead time — short-horizon predictions benefit from recent precipitation and dry-spell features, long-horizon predictions benefit from cyclic phase features and lag lookups.

D. Four-leg multi-track ensemble

The final prediction blends four legs with weakly correlated errors:

- **Anchor (ext150):** a 1071-feature GBDT ensemble (LightGBM + HistGradientBoosting, multiple seeds), followed by the *extrapolate-150* sharpening transform that re-projects predictions toward the empirical 0–5 tails. Public MAE alone: 0.8534.
- **Track 2 (Lag legs):** autoregressive seasonal lookups at $\{1820, 2184, 2215, 2367\}$ days over the historical-label series in `train.csv`, with a light GBDT residual correction. The 2215-day leg targets the dominant ≈ 6 -year ACF peak; its iter4 per-region residual correlation with the GBDT anchor is $\rho=0.51$ (Section IV-G).
- **Track 3 (Dilated CNN):** a 1-D dilated causal convolutional encoder over the 91-day weather sequence, trained with Huber loss and a region embedding.
- **Track 1 (SSL Transformer):** a Time-MAE-style self-supervised Transformer pre-trained on weather-only 91-day sequences and fine-tuned with test-time training (TTT) to mitigate the OOD shift in the `cutoff_age` feature.

Each leg is calibrated to the same mean and standard deviation as the anchor before blending. The blend is

$$\hat{y} = \text{clip}(\alpha \hat{y}_{\text{anchor}} + \beta \hat{y}_{\text{lag}} + \gamma \hat{y}_{\text{cnn}} + \delta \hat{y}_{\text{ssl}}, 0, 5),$$

with weights tuned on a public-like CV slice (Section IV-D). The Greek-letter notation maps to the abstract’s training-menu letters as $(\alpha, \beta, \gamma, \delta) \equiv (A, C, B/2, B/2)$, with the deep-variant term $D=0.10$ folded into the deep family at inference: i.e. $A=0.20$ anchor, $C=0.25$ lag, $B=0.45$ deep (CNN+SSL aggregate), $D=0.10$ deep variants. A final affine post-processor $\hat{y} \leftarrow \text{clip}((\hat{y} - \bar{y}) \cdot s + \bar{y} + \mu, 0, 3)$ with $(\mu, s) \approx (-0.16, 0.98)$ calibrates the submission to the observed public-leaderboard severity plateau ($\bar{y} \approx 1.21$); this public-slice calibration risk is treated as a limitation rather than a generalization claim.

E. Calibration gate

A linear model with features (oracle_mae, MAD, mean, std, high_frac, ρ) maps each candidate to an expected public MAE; the model is fit on 32 historical (candidate, public) pairs from `reports/_local_eval_gate_report.csv`. Only candidates whose predicted public score improves on 0.8534 are uploaded. The fit is assessed by leave-one-out cross-validation (LOOCV): $R_{\text{in-sample}}^2 = 0.9855$, $\text{RMSE}_{\text{in-sample}} = 0.0103$; $R_{\text{LOOCV}}^2 = 0.9472$, $\text{RMSE}_{\text{LOOCV}} = 0.0196$. The 0.038 drop from in-sample

TABLE I: Dataset statistics (transcribed from `reports/data_characteristics_v1.md`).

Quantity	Value
# regions	2,248
# daily rows per region (<code>train.csv</code>)	5,480
# weekly score anchors per region	782
# meteorological feature columns	16
# supervised training samples (after 91-day windowing)	1,485,928
# validation samples (time-based split)	233,792
Public test target mean (all-zero submission MAE)	1.2088
Score support (integer)	{0, 1, 2, 3, 4, 5}
Dominant ACF period of weekly scores	≈ 2184 d (≈ 6.0 yr)
Median test <code>cutoff_age</code> (last train label to test)	≈ 531 d

to LOOCV is modest within this fixed upload ledger, but the doubling of RMSE confirms the gate’s ≈ 0.02 uncertainty band on any single new candidate and does not constitute nested validation of the full search process.

IV. EXPERIMENTS

A. Dataset statistics

Table I summarises the corpus. The synthetic data covers $R = 2,248$ regions over 5,480 daily rows each, giving roughly 1.49×10^6 training samples after the 91-day windowing operation. The score marginal is heavily zero-inflated (the all-zero submission gives public MAE 1.2088, matching the public-leaderboard target mean) and the public-leaderboard slice appears to land in a high-severity phase of the data-generating cycle.

B. Self-defined baselines

Table II reports the baseline ladder. The five median-family baselines come from `reports/baselines.json`; the six OOF MAEs come from `reports/oof_tensor.json`; the two public-leaderboard numbers (ext150 anchor and final blend) come from `reports/v18_morning_report.md`. The ext150 row reports its public MAE because no per-anchor ext150 prediction CSV is preserved on disk.

Two observations: (1) the heuristic “last-observed-score copy” has the best *OOF* MAE of any baseline (0.2369) yet a much worse *public* MAE (1.0884) — a textbook leakage-versus-distribution-shift gap that motivates our public-like CV design; (2) the OOF gap between *b91* and *pl* legs (~ 0.05) confirms that feature engineering (*b91* uses the full 1071-feature pipeline; *pl* is a simpler horizon-pooled variant) matters more than seed choice.

C. Per-horizon MAE breakdown

Table III decomposes each OOF leg by horizon. The MAE rises monotonically from $h=1$ to $h=5$ for every leg, consistent with the information-decay justification for per-horizon modelling.

D. Public submission trajectory and controlled ablation

Table IV reports the public-leaderboard submission trajectory: each row is a single uploaded artefact from the `reports/_local_eval_gate_report.csv` ledger, so the deltas are *not* a controlled held-out ablation. The controlled OOF-aligned ablation that decomposes the

TABLE II: Baseline ladder. “OOF MAE” is on the time-based validation split (233,792 samples). “Public” is the Kaggle public-leaderboard score (40% of test). “–” = artefact not retained on disk.

Method	OOF MAE	Public MAE
<i>Heuristic baselines</i> (<code>reports/baselines.json</code>)		
Last-observed-score copy	0.2369	1.0884
Global median	0.6051	1.2088
Target-month median	0.6051	–
Region median	0.6512	–
Region $\times$ target-month median	0.6691	–
<i>Single learned legs</i> (<code>reports/oof_tensor.json</code>)		
b91 LightGBM (seed 114)	0.3814	–
b91 LightGBM (seed 271828)	0.3753	–
b91 HistGradientBoosting (seed 31415)	0.3792	–
pl LightGBM (seed 114)	0.4312	–
pl LightGBM (seed 271828)	0.4290	–
pl HistGradientBoosting (seed 31415)	0.4253	–
Simple mean of the 6 legs above	0.3910	–
<i>Team milestones</i>		
ext150 anchor (prior team best)	–	0.8534
v18 7-way blend (mid-iteration)	–	0.7952
Final: phd-below075 (this work)	–	0.7628

TABLE III: Per-horizon validation MAE for the 6 OOF legs, computed from `reports/oof_tensor.csv` via `groupby("horizon").apply(MAE)`. “All” is the average across all 33,720 OOF rows.

Leg	All	h1	h2	h3	h4	h5
b91 LGBM s114	0.3814	0.3572	0.3713	0.3915	0.3849	0.4020
b91 LGBM s271828	0.3753	0.3461	0.3654	0.3874	0.3853	0.3924
b91 HGB s31415	0.3792	0.3494	0.3711	0.3862	0.3874	0.4016
pl LGBM s114	0.4312	0.4065	0.4310	0.4436	0.4265	0.4482
pl LGBM s271828	0.4290	0.4079	0.4217	0.4422	0.4305	0.4426
pl HGB s31415	0.4253	0.3951	0.4180	0.4395	0.4257	0.4483
Simple-mean ensemble	0.3910	0.3656	0.3839	0.4023	0.3937	0.4096

same nine method components is Table V; we retain the trajectory because three of its rows have public-leaderboard ground truth that the OOF tensor cannot supply. OOF MAE in Table V is used to audit relative component behaviour across the menu, not as a predictor of public or private leaderboard score; the OOF-vs-public mismatch documented in Table II (e.g., last-observed-score copy has OOF 0.2369 but public 1.0884) is a leakage-versus-distribution-shift gap, not a contradiction.

E. Multi-year periodicity

Figure 1 reports, from real data, the multi-year periodicity that the lag legs exploit. Panel (a) shows the distribution of the dominant FFT period across all 2,248 regions: a clear mass appears in the 5–8 year band (highlighted in red), and a separate mass at ≈ 15 years corresponds to a harmonic of the same underlying cycle. Panel (b) plots the raw weekly score series for two example regions, which visibly oscillate on a multi-year envelope. Panel (c) quantifies the per-region distribution of the dominant period: $\sim 30\%$ of regions fall in the 5–8 year band that our lag-2215 leg targets, and a similar fraction sits at the ≥ 12 -year harmonic. The panel does *not* claim a per-region ACF at exactly 2215 days; the recomputed iter4

TABLE IV: Public submission trajectory — chronological record of uploaded public submissions, not a controlled validation-set ablation. “ Δ ” is the public-MAE change vs. the row above. The aligned-OOF controlled ablation that decomposes the same component stack is Table V. Configurations are sourced from `reports/_local_eval_gate_report.csv`.

Configuration	Public MAE	Δ
ext150 anchor (1071-feat GBDT + extrapolate-150)	0.8534	—
+ 3-way blend with lag legs	0.8100	−0.0434
+ 4-way (lag + Track 3 CNN)	0.8017	−0.0083
+ 7-way (add SSL Transformer + harmonics)	0.7952	−0.0065
+ affine post-processor + clip 3.0 (final)	0.7628	−0.0324

TABLE V: Controlled OOF-aligned 9-row ablation (iter4). All rows are evaluated on the same 33,720-row OOF tensor (6,744 anchors $\times$ 5 horizons) produced by `scripts/build_ablation_9row.py`. “Public MAE” is filled when an uploaded checkpoint matches the configuration; “—” otherwise. The opposite ordering of OOF vs. public MAE between rows 8 (no affine) and 9 (affine+clip, the final) is direct evidence that the affine-and-clip step is public-slice calibration — it trades OOF accuracy for proximity to the public severity plateau (Threat (iv), Section IV-I).

Configuration	OOF MAE	Public MAE
1. Weather-only single GBDT (best of 6 legs)	0.3753	—
2. Lag-2215 baseline (autoregressive lookup)	0.9673	—
3. Lag-2215 + per-horizon bias correction	1.0694	—
4. GBDT anchor (6-leg b91+pl blend)	0.3910	0.8534
5. GBDT + lag (A+C renormalised, $0.44A + 0.56C$)	0.6840	—
6. GBDT + lag + CNN (Track 3)	0.5389	0.8017
7. GBDT + lag + CNN + SSL Transformer	0.5523	0.7952
8. Full blend, no affine/clip ($0.2A + 0.45B + 0.25C + 0.1D$)	0.5325	—
9. Full blend + affine + clip 3.0 (final)	0.5766	0.7628

per-region residual correlations between non-GBDT legs and the GBDT anchor are reported in Section IV-G.

F. Observed public-ledger slope trend

Figure 2 shows, from the 32 historical (MAD, public) uploads in `reports/_local_eval_gate_report.csv`, the empirical positive-slope relationship between MAD-from-anchor and public MAE. Blue circles are pre-v18 candidates; orange triangles are the 3 v18+ candidates. A submission-row bootstrap fit ($B=1000$) over all 32 points gives slope 0.20 (95% CI [0.05, 0.26]) and intercept 0.85 (95% CI [0.82, 0.88]). Because these rows are adaptive submissions rather than independent experimental units, we treat the positive slope as descriptive evidence of an observed public-leaderboard trend, not as a standalone statistical law. The v18+ orange cluster is the first group in this ledger to drop below the bootstrap-fitted trend line.

The bootstrap restricted to the 29 pre-v18 rows only also gives a slope near 0.21; the often-quoted earlier number of ≈ 0.42 from internal team analyses was an upper-tail subset, not the population fit. We present the bootstrap-fitted line in Fig. 2 from real data (rather than the older anecdotal 0.42) because it is the auditable estimate




Fig. 1: Multi-year periodicity in the weekly drought scores — evidence for the lag-leg design choice. (a) Histogram of dominant FFT period per region across all $R=2,248$ regions (in years; 5–8-year band shaded). (b) Example weekly score series for two regions (R1 and R1001) showing visible multi-year structure. (c) Per-region distribution of the dominant FFT period, binned into 5 ranges. Data source: `reports/_track2_fft_peaks.csv` and `data/train.csv`.

over the actual evidence in `reports/_local_eval_gate_report.csv`.

G. Cross-family residual diversity (real OOF)

Figure 3 is the real 6×6 region-mean residual correlation matrix on the 33,720-row OOF tensor (`reports/oof_tensor.csv`). Two clean blocks are visible: within the b91 family ($\rho \approx 0.98$ across LGBM/HGB seeds) and within the pl family ($\rho \approx 0.98$), with the cross-family value only $\rho \approx 0.79$ – 0.82 . This is the directly reproducible justification for keeping both tabular feature families in the ensemble.

For the cross-family residual correlation against the GBDT anchor (the lag-2215, CNN, and SSL legs vs. the 6-leg anchor), iter4 regenerates the missing OOF predictions for the lag-2215 lookup (`scripts/regen_lag_2215_oof.py`, deterministic over `train.csv`; runtime ≈ 9 s) and for the SSL Transformer (`scripts/regen_ssl_oof.py`, single-pass forward of `checkpoints/track1_finetuned.pt` on the same




Fig. 2: Observed slope trend on real submission data. Each marker is one Kaggle upload from the 32-row `reports/_local_eval_gate_report.csv` gate ledger. Dashed line: submission-row bootstrap-fitted slope on all 32 rows (0.20, 95% CI [0.05, 0.26], shaded band B=1000). Dotted line: bootstrap-fitted slope on the 29 pre-v18 rows only (0.21). Solid grey line: ext150 ceiling 0.8534. v18+ candidates (orange triangles) are the first uploads in this ledger to drop below the bootstrap-fitted line, with the lowest sitting at 0.7952. The final affine-and-clip 0.7628 submission is excluded from this slope plot because affine post-processing makes the MAD axis non-comparable; see Fig. 5 and Table IV for the full trajectory.

6,744 validation anchors; runtime ≈ 1 s on GPU). Per-region Pearson correlation between $r_{\text{leg}} = y - \hat{y}_{\text{leg}}$ and $r_{\text{anchor}} = y - \hat{y}_{\text{anchor}}$ is then resampled across regions (B=1000, `scripts/compute_cross_leg_rho.py`, following Efron [7]):

- lag-2215: mean $\rho = 0.51$, 95% CI [0.485, 0.533] ($n=1,374$ regions with ≥ 3 valid pairs);
- CNN (mean of 5 cached Track-3 runs): mean $\rho = 0.54$, 95% CI [0.524, 0.562] ($n=2,248$);
- SSL Transformer: mean $\rho = 0.33$, 95% CI [0.311, 0.349] ($n=2,248$).

These values reflect *cross-family residual diversity* rather than strict orthogonality: SSL is materially lower-correlated with the GBDT anchor than CNN or lag-2215, but $\rho \approx 0.33$ does not constitute independence. The earlier training-menu figure “ $\rho = 0.107 \pm 0.027$ ” for lag-2215 (from `reports/data_characteristics_v1.md`) was computed under a different aggregation that the original artefacts no longer permit auditing; the iter4 recomputation above supersedes it. Cached file SHA256 hashes are in `ARTIFACTS.md`.




Fig. 3: Real 6×6 region-mean residual correlation matrix for the 6 OOF GBDT legs (2,248 regions, computed from `reports/oof_tensor.csv`). The two ≈ 0.98 within-family blocks correspond to the *b91* and *pl* feature families; the off-diagonal ≈ 0.80 value is the within-tabular ensemble diversity. Recomputed cross-family residual correlations of the lag-2215, CNN, and SSL Transformer legs against the 6-leg GBDT anchor are reported in Section IV-G.

H. 5-fold region CV and trajectory

Figures 4 and 5 support the 5-fold region CV and submission-trajectory claims.

I. Threats to validity

We list four threats and the evidence that bounds each. (i) *Calibration-gate overfit*. The 32-row gate fit has $R^2_{\text{in-sample}} = 0.9855$ but $R^2_{\text{LOOCV}} = 0.9472$; the LOO RMSE of 0.0196 is the uncertainty band we attach to any one new candidate’s predicted public MAE.

(ii) *Slope-trend point estimate*. The submission-row bootstrap CI on the slope (B=1000) excludes zero ([0.05, 0.26]), supporting a positive association in the 32-upload ledger. Because the uploads were adaptively chosen, this is descriptive evidence rather than a controlled statistical test.

(iii) *Cross-family residual diversity (iter4 update)*. The iter4 recomputation supersedes the earlier $\rho = 0.107 \pm 0.027$ figure from `reports/data_characteristics_v1.md`: the regenerated and row-aligned OOF predictions for the lag-2215 leg and the SSL Transformer now allow a direct, bootstrap-CI bounded per-region Pearson computation (Section IV-G). The new values are $\rho=0.51$ (lag-2215), 0.54 (CNN), 0.33 (SSL). These are *lower-correlated* rather than orthogonal; SSL is the materially lowest. The lower SSL

value most plausibly explains the SSL leg’s positive contribution in Table V; lag-2215 contributes primarily through label-distribution coverage rather than residual independence.

(iv) *Public/private regime shift.* The 0.7628 result was achieved by tuning on public-test-like CV folds (high-severity phase, target mean ≈ 1.21). The training menu’s convex weights are $A=0.20$ (GBDT anchor) + $B=0.45$ (deep CNN+SSL) + $C=0.25$ (lag) + $D=0.10$ (deep variants); the lag component is therefore only 25% of the blend, and if it attenuates on a private slice landing in a different phase of the 6-year cycle the remaining 75% from weather-driven components ($A+B+D$) still covers the prediction. We do not state a numeric private-LB prediction. Direct evidence that the final affine-and-clip step is a *public-slice calibration* rather than a general-purpose improvement comes from Table V: the full blend without affine/clip (row 8) has OOF MAE 0.5325, while the affine-and-clip variant (row 9) has OOF MAE 0.5766 — worse on the OOF tensor — yet wins on public (0.7628). The step trades OOF accuracy for proximity to the public severity plateau and is therefore a competition-calibration step with explicit private-LB risk. If Kaggle’s automatic two-submission selection allows it, the no-affine 7-way row of Table IV (public MAE 0.7952) is the natural private-LB-conservative sibling to the 0.7628 final.

V. REPRODUCIBILITY

A. Environment

Python 3.11; the relevant pinned versions are `numpy` 1.26, `pandas` 2.1, `scikit-learn` 1.4, `lightgbm` 4.3, `matplotlib` 3.8. The complete environment is specified in the repository’s `requirements.txt`.

B. Single-command cached re-blend of the 0.7628 submission

The uploaded artefact basename `submission_phd_below075_20260522.csv` lives under `submissions/`. It is regenerated from cached predictions and the training menu by `make phd-below075`, which expands to the two commands shown in Figure 6. This target does not retrain every historical model from raw data; it reuses the retained `v18/cache` artefacts documented in `ARTIFACTS.md`. The regenerated CSV matches the uploaded artefact with maximum absolute difference below 5×10^{-16} in `make verify-submission`; all numerical comparisons in this report are stable at the 4-decimal precision used throughout.

C. Build of this report

`cd report && make` runs the figure generator (`report/figures/generate_figures.py`) and two passes of `pdflatex` to produce the IEEE A4 PDF. The Makefile also exposes a `make check` target that verifies the headline strings (Team 3, 0.7628, GitHub URL) appear in the PDF.

VI. LIMITATIONS

Synthetic data. The competition data is a synthetic re-rendering of a US-county-level drought corpus; our model’s improvements may not transfer linearly to a true operational drought monitor.

Public/private regime risk. The 0.7628 result was achieved by tuning on a public-leaderboard-like CV phase. If the private-leaderboard slice lands in a different phase of the 6-year cycle, the lag-leg contribution may attenuate; the GBDT anchor would then carry more of the prediction. The final affine shift and clip upper bound are especially public-slice-sensitive, so we do not state a numeric private-leaderboard prediction. **We interpret 0.7628 as a public-leaderboard result, not as an unbiased estimate of private test MAE**; the affine-and-clip step is in particular a public-slice calibration (see Threats (iv) in Section IV-I and the OOF row 9 vs. row 8 inversion in Table V).



OOD on cutoff_age. The test gap between the last training label and the prediction window has median ≈ 531 days, outside the deepest training densities; Track 3’s test-time training is an explicit mitigation.

Calibration gate sample size. The 32-point linear gate is small (see Threats to Validity in Section IV-I); its LOOCV $R^2 = 0.9472$ implies a ± 0.02 uncertainty band on any one new candidate’s predicted public MAE.


Computational budget. The cached re-blend/report path runs in roughly 1.5 hours on a single 16-core CPU node plus one consumer GPU (RTX 4090 for Track 3 cache generation); the training menu and OOF tensors total ≈ 4.5 GB on disk.

What we did not measure. We did not retrain the GBDT anchor, the lag legs, or the deep tracks for this report — the goal is reproducible description of the cached artefact that scored 0.7628, not a fresh hyperparameter sweep. Per-region failure-mode analysis (which regions does the blend hurt rather than help?) is left for future work.

(a) Calibration OOF

(b) Lag- ρ Stability

(c) Blend Weight MAE Robustness



(d) Final Candidate Stability




Fig. 4: 5-fold region CV evidence (panels reproduced from `reports/plots/cv1...cv4*.png`). (a) Calibration OOF scatter (RMSE 0.015 between in-sample and OOF predictions); (b) lag- ρ stability across folds as recorded in the original CV run (the underlying $\rho \approx 0.107$ figure is superseded by the iter4 recomputation in Section IV-G; the panel is retained as historical context for the design choice); (c) in-sample vs. OOF blend MAE ratio ≈ 1.001 for the chosen weights; (d) per-fold final-candidate MAE for the 0.7628 blend. These diagnostics suggest the chosen weights are not driven by a single fold, but they are not a substitute for private-leaderboard validation.




Fig. 5: Submission trajectory on the public leaderboard. From the prior team ceiling of 0.8534 (ext150) through the 0.7952 mid-iteration blend to the final 0.7628 phd-below075 artefact — a 10.6% relative improvement.

```
PYTHONPATH=src python \
scripts/analyze_data_distribution.py \
--force-synthesis --emit-menu

PYTHONPATH=src python \
scripts/multi_blend_grid.py \
--menu reports/training_menu_v1.json \
--fixed \
--out submissions/submission_\
phd_below075_20260522.csv
```

Fig. 6: The two commands executed by `make phd-below075`. This is the cached re-blend path used to reproduce the uploaded CSV; `make verify-submission` compares the output against the archived artefact. The `--force-synthesis` flag triggers training-menu re-emission from cached per-region features; it does not synthesise labels, predictions, or submissions (see the repository `README.md` flag glossary).

Author Contributions. Team 3 (DM 114, Spring 2026) members: 劉宇舜 (413551030), 張東浩 (413551038), 陳冠甫 (414551017). Specific contributions are recorded in the GitHub repository’s contribution log; the team collectively assumes responsibility for the design, the reported numbers, and the reproducibility of the artefact described above. The report itself was assembled iteratively from internal experiment logs (`docs/JOURNEY.md`, `reports/v18_morning_report.md`) and the artefact ledger (`reports/_local_eval_gate_report.csv`, `reports/oof_tensor.csv`); statistical claims were audited by leave-one-out cross-validation and bootstrap resampling against those artefacts before inclusion.

REFERENCES

- [1] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A highly efficient gradient boosting decision tree,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.
- [2] A. van den Oord, S. Dieleman, H. Zen, K. Simonyan, O. Vinyals, A. Graves, N. Kalchbrenner, A. Senior, and K. Kavukcuoglu, “WaveNet: A generative model for raw audio,” *arXiv preprint arXiv:1609.03499*, 2016.
- [3] G. Kechyn, L. Yu, Y. Zang, and S. Kechyn, “Sales forecasting using WaveNet within the framework of the Kaggle competition,” *arXiv preprint arXiv:1803.04037*, 2018.
- [4] A. F. Ansari, L. Stella, C. Turkmen, X. Zhang, P. Mercado, H. Shen, O. Shchur, S. S. Rangapuram, S. P. Arango, S. Kapoor, J. Zschiegner, D. C. Maddix, M. W. Mahoney, K. Torkkola, A. G. Wilson, M. Bohlke-Schneider, and Y. Wang, “Chronos: Learning the language of time series,” *arXiv preprint arXiv:2403.07815*, 2024.
- [5] B. Lim, S. Ö. Arik, N. Loeff, and T. Pfister, “Temporal fusion transformers for interpretable multi-horizon time series forecasting,” *International Journal of Forecasting*, vol. 37, no. 4, pp. 1748–1764, 2021.
- [6] M. Svoboda, D. LeComte, M. Hayes, R. Heim, K. Gleason, J. Angel, B. Rippey, R. Tinker, M. Palecki, D. Stooksbury, D. Miskus, and S. Stephens, “The drought monitor,” *Bulletin of the American Meteorological Society*, vol. 83, no. 8, pp. 1181–1190, 2002.
- [7] B. Efron, “Bootstrap methods: Another look at the jackknife,” *Annals of Statistics*, vol. 7, no. 1, pp. 1–26, 1979.